

When Not to Generate: How AI Systems Quote Scripture, and What Authoritative Quotation Requires

Alex Chao
Fide AI
alex@fideai.org

July 2026

Abstract

People increasingly ask AI assistants for Scripture—in sermon preparation, teaching, and pastoral conversation. These are not ordinary information requests: someone asking for a passage wants particular words, in a particular translation, from a source they could have opened themselves. A reply that blends two translations, drops a verse, or paraphrases inside quotation marks satisfies none of that while appearing to satisfy all of it, and the person who asked is the one least able to check. Existing evaluation does not see this failure. A response can be factually accurate, properly cited, and semantically faithful, and still be the wrong text.

When the words themselves are the answer, who should be responsible for them? Locked in advance, we route 8,640 matched Scripture requests through four designs across six model families, three open English editions, and twenty passages. Quoting from memory delivered the requested edition exactly 25.00% of the time, an authorized lookup tool 80.09%, and deterministic insertion 83.94%. Supplying the passage directly reached 93.61%, but that condition is an oracle rather than a rival: it never selects or retrieves, and no deployed system knows the answer in advance. Connecting a source did not remove failure; it moved failure onto whatever the model still decided. Models called the tool in 95.00% of cases yet asked it for the right passage in only 84.17%. Holding the source constant sharpens the practical finding: handed the exact text, a generative layer still altered it 6.39% of the time, against 0.11% for non-generative rendering. Exact quotation is not one skill but a chain, and each link fails on its own terms. This measures textual delivery, not whether a passage was well chosen, rightly interpreted, or pastorally appropriate.

1 Introduction

Consider three requests a person might make of an AI assistant. “Give me Ephesians 2:8–9 in the Berean Standard Bible for Sunday’s slide.” “Quote the passage where Jesus compares his being lifted up to Moses’ serpent in the wilderness.” “Send me Revelation 21:1–4 to read at a funeral.” Each asks for exact words, but each puts a different decision in the system’s hands. The first names the passage and the edition, leaving only faithful reproduction. The second names neither, so the passage must be identified before it can be quoted. The third is long enough that truncation, not misidentification, is the likely failure. An assistant that answers all three the same way—by generating fluent text that carries the right reference—fails each of them differently, and the final prose does not reveal which failure occurred.

That opacity is the practical problem. The meaning may survive a blended or shortened rendering; the quotation does not. And the person who asked is poorly placed to notice, because checking the answer would require already having the text they asked for. For requests like these, wording is not merely a vehicle for the answer. Wording is part of the answer.

These requests are not hypothetical, and the settings are not peripheral. In a December 2025 survey of U.S. Protestant pastors, 87% reported using AI in ministry, 36% for biblical or theological research, and 24% for writing or editing sermons (Barna Group, 2026). A separate survey of 594 church leaders found comparable adoption alongside near-absent governance: most reported frequent use, while 73% reported no AI policy of any kind (Exponential AI NEXT and ChurchTechToday, 2025). Which English translation a person reads is itself a durable, measurable feature of American religious life rather than an incidental preference (Goff et al., 2017), so substituting one edition for another is not a cosmetic error.

The stakes are practical and specific. A congregation reads responsively from a printed edition while a screen displays different wording. A student memorizes a verse in one translation and is later corrected against another. A teacher builds an argument on a clause that the quoted edition does not contain. Someone in distress receives a passage with its final line missing. In each case the response is fluent, carries the right reference, and is wrong in the way that matters. Readers reasonably treat quotation marks and a verse reference as a claim about a source of record; when that claim is unreliable and difficult to detect, the cost falls on the reader rather than on the system.

Practitioners have noticed. The chief executive of a major digital Bible platform has publicly estimated that leading models misquote Scripture between 15% and 60% of the time, and has cited that estimate as a reason not to ship a Scripture-answering assistant (Christian Daily International, 2026). The estimate is plausible and consequential, but no measurement protocol accompanies it: the systems, passages, editions, and criterion for “misquote” are unstated, so it can be neither checked nor acted on with confidence. Claims of that importance deserve a specified, reproducible measurement. Providing one, for a bounded set of English editions and conditions, is the purpose of this paper.

We use *authoritative quotation* to mean delivering the requested span, in the requested edition, from an identified source of record—a specific, versioned text designated as authoritative for that request, rather than an abstract notion of authority. A response may be relevant, factually close, or theologically unobjectionable while still failing that narrower task. It may silently substitute an edition, blend familiar formulations, omit a verse, expand beyond the requested span, or present a paraphrase as a quotation. This is a stricter requirement than factuality, attribution, or groundedness: a response can be factually supported and properly cited while still failing an exact quotation request.

Neither model scale nor retrieval resolves the problem on its own. Language models are built to generate plausible continuations; parametric memory can sometimes reproduce a familiar source verbatim, but it is not a source-of-record interface. Retrieval can make authoritative text available, and a generative layer can still alter it. A model holding a tool can bypass the tool or request the wrong span. A deterministic renderer can preserve source bytes exactly, but only downstream of a correct decision about source, edition, and boundaries. Return to the serpent request above and the point sharpens: a model answering from memory may blend familiar translations; a model given the source may shorten it; a tool-enabled model may retrieve the neighboring verse; and a deterministic renderer may reproduce the source perfectly after the model selected the wrong range. The four designs fail at four different places, and every one of them produces credible prose. Authoritative quotation is a property of the whole source-delivery chain, not a model knowledge property.

We therefore hold the user’s request fixed and vary which component owns the words. Matched

requests were routed through four architectures: generation from model memory, reproduction of source text supplied in context, retrieval through an authorized `get_passage` tool, and model selection of a structured reference followed by deterministic rendering. The prospectively specified and internally hash-locked design contains 8,640 observations across six model families, three open English editions, twenty passages, two prompt forms, and three repetitions. This is a matched crossed comparison rather than a benchmark or leaderboard, and the primary endpoint is deliberately demanding: the user-visible output must equal the requested text exactly, with condition-specific evidence that the declared architecture was actually followed. Scripture makes this comparison unusually tractable—verse references supply a conventional addressing scheme, multiple English translations make edition identity observable rather than assumed, span length can be varied systematically, and indirect prompts force semantic identification before rendering.

The results are stark and, we think, more interesting than a ranking. Native generation from model memory met the strict endpoint in 25.00% of observations, against 93.61% for source-supplied quotation, 80.09% for tool-mediated retrieval, and 83.94% for deterministic rendering. Because the endpoint requires the *requested* edition, much of the native gap is edition substitution rather than general unfamiliarity: native exactness ranged from 41.53% for the Berean Standard Bible to 5.00% for the Literal Standard Version. The central finding is what happened after a source was connected. Failure did not disappear; it migrated toward whichever decision the architecture still left to the model. A source tool was invoked in 95.00% of tool observations but used for the requested lookup in only 84.17%. Deterministic rendering was exact in 1,813 of 1,815 cases once the model had selected the correct reference—but the model selected correctly in only 84.03% of observations, and no renderer can repair a wrong selection. Each architecture made one stage reliable and left another exposed.

What transfers from this is the structure, not the numbers. The same shape appears wherever a source of record exists: a dosing instruction, a statutory subsection, an indemnity clause, and a safety standard each have an edition, a boundary, and wording that a paraphrase does not preserve, and each can fail at selection, access, rendering, or final output. The decomposition and evaluation design developed here apply to those settings directly. The measured rates do not—every number in this paper is about English Bible quotation, and another domain needs its own sources, scenarios, and expertise. Many faith traditions likewise designate specific textual forms as authoritative and treat exact wording as consequential, but the unit of authority differs—a source-language consonantal text, a recited form, a received translation, a commentary tradition—and those differences would change the study design, not only the results. We report on English Christian Bible editions and make no claim about any other tradition’s texts. Textual fidelity is also not theological fidelity: the study evaluates whether a system delivered the requested source faithfully, not whether the passage was well chosen, rightly interpreted, or pastorally appropriate.

Contributions. The paper makes three contributions. First, it gives a rigorous empirical account of Scripture quotation across twenty passages, three open English editions, six model families, and four system architectures. Second, it defines a *source-delivery chain* that separates source selection, authorized access, exact rendering, and final-output integrity, and shows that quotation failures are only interpretable when attributed to a stage. Third, it releases an open evaluation methodology, deidentified derived scores, and reproducible analyses while keeping private execution code, raw model traces, and rights-constrained text outside the public artifact. The study answers a public call for research on sacred-text fidelity (Fide AI, 2026), and the repository records which parts of that call remain open. Three questions organize the study: how architecture-adherent exactness differs across the four architectures; at which stage failures occur; and how prompt form, passage

length, edition, model family, and repeated sampling interact with those failures.

2 Authoritative Quotation as a Systems Problem

2.1 Quotation succeeds only when the whole chain succeeds

A successful quotation needs four things to go right. The system must identify what text the user means, obtain it through the declared path, preserve its wording, and deliver that wording intact. Failure at any one stage defeats the quotation request even if every other stage succeeds.

Formally, let a request specify, explicitly or implicitly, a source work W , an edition E , and a span B . Let $T(W, E, B)$ be the authoritative text returned by the source of record. We separate four system properties:

- S : the intended work, edition, and span are selected correctly;
- A : the declared source-access method is followed and the requested source is available;
- R : the selected source text is rendered without alteration; and
- I : the user-visible final output preserves that rendering without truncation, leakage, or extraneous text.

The strict endpoint is conjunctive:

$$\text{E2E-EXACT} = S \wedge A \wedge R \wedge I. \quad (1)$$

The operational evidence for each term depends on architecture. In native generation there is no external access event to verify; exact final output is the decisive behavior. In a tool condition, access requires an observed call for the requested span rather than merely the presence of a tool definition. In a deterministic condition, exact rendering is meaningful only after a parseable, correct reference has reached the renderer.

This decomposition prevents a common category error. A renderer may provide a strong guarantee about R while the overall system still fails because S is false. Conversely, a model may identify the correct passage while a generative surface layer changes the words. An aggregate score is useful, but stage-level evidence is required to explain it.

2.2 Why factuality and citation are not enough

A quotation can be factually supported yet textually wrong. A citation can point to the correct document while the quoted span omits a limiting clause. A semantic similarity score can treat two Bible translations as nearly equivalent even when the user requested one exact edition. Authoritative quotation is therefore stricter than factuality, attribution, or general groundedness.

The requirement is also narrower. The evaluation does not ask whether the source itself is true, wise, clinically appropriate, legally controlling, or theologically interpreted correctly. It asks whether a system delivered the requested source faithfully under its declared architecture.

2.3 Scripture as the study domain

Scripture warrants study because AI systems increasingly mediate access to texts that Christian communities receive, interpret, and transmit as authoritative, and because the domain makes the quotation problem measurable. A reference such as “John 3:16” identifies a conventional verse

span but not one English wording, so edition identity is observable rather than assumed. An event-style request adds a prior selection problem: the system must identify the intended span before it can quote it. Few domains combine a stable addressing convention, multiple simultaneously authoritative editions, and a community that treats exact wording as consequential.

Working here requires textual and theological restraint. The English editions studied sit within a complex manuscript and translation history in both the Hebrew Bible and New Testament (Tov, 2012; Metzger and Ehrman, 2005), and translation requires linguistic and interpretive decisions (Nida and Taber, 1969). The three editions differ not only in translation philosophy but in underlying text-critical basis, so “edition” is partly a claim about source-language witnesses rather than purely a matter of English wording. The twenty targets were not selected to probe variant units, and a study designed around textual variants would be a distinct and worthwhile contribution. Exact reproduction of a selected English edition is therefore not fidelity to the Hebrew, Aramaic, or Greek source texts, and no result here establishes that one English translation is doctrinally or linguistically superior. The twenty targets all fall within books shared by the major Christian canons; the study does not adjudicate canonical boundaries, and the three editions carry the Protestant canon.

For churches and Christian institutions, the practical implication is correspondingly bounded. A system that quotes the requested edition exactly removes one avoidable textual failure. It does not replace biblical scholarship, ecclesial authority, pastoral judgment, or attention to literary and canonical context. Exact quotation is a prerequisite for some forms of responsible use, not a complete account of responsible use.

3 Related Work

Retrieval-augmented generation combines parametric sequence models with retrieved evidence (Lewis et al., 2020). Later work examines when parametric memory is insufficient and non-parametric memory should be preferred (Mallen et al., 2023), while RAGAs evaluates retrieval and generation quality within retrieval-augmented systems (Es et al., 2024). Tool-use research asks whether models can decide when and how to invoke external APIs (Schick et al., 2023; Ross et al., 2025), and grammar-constrained decoding offers one way to enforce structured interfaces (Geng et al., 2023). Authoritative quotation joins these concerns: the system must obtain a declared span, demonstrate that the intended access path was used, and preserve the result.

Fine-grained factuality work decomposes long-form generation into checkable claims (Min et al., 2023), while citation-oriented systems evaluate whether generated claims are accompanied by supporting sources (Gao et al., 2023). Surveys document broad forms of unsupported and unfaithful neural generation (Ji et al., 2023; Huang et al., 2025). Our task is stricter in wording and narrower in scope: a response can be factually supported and properly cited while still failing an exact quotation request.

CopyBench measures literal and non-literal reproduction of copyrighted fiction books (Chen et al., 2024), while Quote-Tuning increases verifiability by aligning models to quote verbatim strings found in trusted pretraining corpora (Zhang et al., 2025). The present study differs in object and architecture: the requested edition and span are specified as source-of-record requirements, and authorized retrieval or deterministic rendering can replace parametric reproduction. These works are complementary and bound the novelty claim; this paper does not claim to introduce verbatim-copy measurement or quoting as a general research problem. Training-data extraction and memorization studies likewise show that models can reproduce memorized content under some conditions (Carlini et al., 2021, 2023). Memorization does not make parametric recall a stable or

authorized source interface. Copyright policy work, including the United States Copyright Office’s generative-AI training report, provides further context for making source access and redistribution assumptions explicit ([United States Copyright Office, 2025](#)).

Holistic and behavioral evaluation emphasizes scenario coverage, multiple metrics, and tests of specific capabilities rather than one undifferentiated score ([Liang et al., 2023](#); [Ribeiro et al., 2020](#)). We adopt that orientation without turning the study into a ranking exercise. The contribution is the source-delivery chain: a matched evaluation of which component selected, obtained, rendered, and preserved the requested source.

4 Study Design

The experiment holds the request panel fixed and changes the route by which a quotation reaches the user. This makes architecture, rather than passage choice or model mix, the primary comparison.

4.1 Prospective specification and execution boundary

The primary study was prospectively specified and internally hash-locked before confirmatory execution began. This was not a public registry deposit. The lock fixed the target pack, source registry, prompt families, model routes, conditions, repetitions, scoring rules, denominator policy, and bootstrap procedure. Integration and calibration runs were excluded from confirmatory estimates. Two post-lock changes were administrative—they removed budget and wall-clock stopping ceilings and split execution into resumable blocks—and changed no prompt, source, model, observation, score, or analysis.

The replication package releases the lock receipt and deviation record together with the executed prompts, tool schema, and sampling policy, so the declared design can be checked against what ran. Authoritative passage text is withheld; its release-safe digest is published in its place.

The primary phase contained:

$$20 \text{ targets} \times 2 \text{ prompt families} \times 4 \text{ conditions} \times 3 \text{ editions} \times 6 \text{ model routes} \times 3 \text{ epochs} = 8,640. \quad (2)$$

The unit of analysis is one target–prompt–condition–edition–model–epoch observation. Conditions are paired within all other factors.

4.2 Passage targets

The target set is a stratified purposive panel, not a probability sample of the Bible or of user requests. It contains five targets in each of four strata:

1. explicit single verses;
2. explicit short multi-verse passages;
3. longer passages intended to expose truncation, span, or guardrail behavior; and
4. event-style or semantic reference requests.

Targets include narrative, poetry, wisdom, prophetic, Gospel, and epistolary material. Two known intertextual or adjacent-passage relationships were annotated rather than treated as fully independent evidence: the Shema and great-commandment pair, and adjacent passages in John 3.

Each target had an explicit-reference prompt and an edition-neutral contextual description. The distinction is who supplies the reference. In the explicit-reference family the caller names it—“Quote [reference] from the [edition] exactly”—so no architecture has to identify the passage, and the comparison isolates obtaining and delivering it. In the contextual family the caller describes the passage without naming book, chapter, or verse, so identification becomes part of the task. Every target was evaluated under both, crossed with all four architectures. The contextual descriptions did not name the book, chapter, or verse and were designed to avoid distinctive wording from any tested edition. An initial review identified six descriptions that needed revision. The revised six were then reviewed again under blinding, and all six were matched to the intended reference. Because the original twenty-item review was not fully blinded, only this six-item result is reported as blinded evidence.

4.3 Source editions

The full four-condition comparison used three open English editions:

- Berean Standard Bible (BSB), a modern public-domain edition ([Berean Bible, 2023](#));
- World English Bible Updated (WEBU), a public-domain translation ([eBible.org, 2024](#)); and
- Literal Standard Version (LSV, 2020), released under CC BY-SA ([Covenant Press, 2020](#)).

These editions differ in translation philosophy and in underlying text-critical basis, so a requested edition specifies more than a style of English.

A source of record is a versioned artifact, so edition identity is disclosed rather than left to the edition name. The public package includes a release-safe source-edition record giving, for each edition, the provider and provider-side source identifier, the edition designation, the rights basis, and a normalized-text SHA-256 digest for each of the twenty targets. Ground truth was verified at 2026-07-26T20:38:10Z by two uncached retrievals through the pinned shared runtime with matching digests; this establishes retrieval stability, not independence of the upstream provider. The same stored text then served as ground truth for every observation. Because the digests are computed over whitespace-collapsed UTF-8 text, a replicator can hash their own copy of a passage and confirm they hold the same edition text without either party transmitting it. The record is derived from the source registry fixed by the prospective lock and carries that artifact’s digest, so the disclosure can be checked against a commitment made before confirmatory execution.

Authoritative source text was retained privately for scoring. The public result package contains aggregate metrics and no LSV passage text. Four restricted editions were present in the locked registry but excluded from the primary phase pending source and source-to-model authorization; they are listed in the source-edition record with that status.

4.4 Model routes and sampling


The primary panel used one route from each of six model families: GPT-5.6 Sol, Claude Sonnet 5, Gemini 3.5 Flash, DeepSeek V4 Pro, Kimi K3, and GLM-5.2. These labels identify the end-points evaluated on July 26–27, 2026; they are not claims about later versions or general vendor performance.

Requests used OpenRouter-compatible routing with one upstream provider pinned, fallbacks disabled, required parameter support, and data-collecting routes denied. Provider-default sampling was used with temperature omitted. Every cell was observed in three epochs. A failed observation was never re-prompted in an attempt to obtain a better answer: semantic retries were zero. Bounded

transport retries were allowed, and terminal provider errors remained non-deliveries in the primary denominator.

4.5 Architecture conditions

The four conditions give the model progressively different responsibilities. In the first, the model must supply the words itself. In the second, the words are already present and the model must preserve them. In the third, the model must obtain the words through a tool and then return them. In the fourth, the model selects only the reference; a deterministic layer owns the words. The caller asks the same question in every matched cell, so the comparison changes the delivery path rather than the user’s request (Figure 1).



The user request is held fixed. What changes is which component owns selection, source access, and the final words.

Figure 1: The four evaluated paths through the source-delivery chain. The caller request is matched across conditions; what changes is which component owns the source words and the decisions that precede them.

Native parametric quotation. The model received no source text and no tool. This condition represents the common behavior of asking a general assistant to quote from memory.

Source-supplied quotation (oracle). The authoritative span was placed in the model context and the model was asked to return it unchanged. This is deliberately not a retrieval architecture: there is no lookup, and no decision about what to retrieve, because the correct span is supplied in advance. It is an oracle condition that isolates a single question—once authoritative text is already available, does the generative layer preserve it?—and it bounds what any perfect retrieval system could deliver. It is reported alongside the three deployable architectures, but it is not one of them, and comparisons that treat it as a peer will overstate it by the amount that source selection costs the others.

Tool-mediated retrieval. The model received an authorized `get_passage` tool and was instructed to use it. The trace recorded whether the tool was invoked and whether it retrieved the

requested span. Emitting text that resembled a tool call without an executable call did not count as tool use.

Structured selection and deterministic rendering. The model was instructed to return exactly one placeholder of the form `<quote>{{QUOTE:<reference>}}</quote>`. A local deterministic layer parsed the reference, fetched the selected edition and span, and replaced the placeholder. The primary score required the strict placeholder grammar exactly as executed. A tolerant parser, specified in advance and used only for analysis, separately assessed whether a single unambiguous reference could still be recovered from malformed output; it never altered the observed output or the primary score.

4.6 Outcomes and estimands

Primary and common outcomes. The locked primary endpoint, E2E-EXACT, required strict final-output equality and all condition-specific selection, lookup, replacement, and method-adherence checks. We call it *architecture-adherent exactness* throughout: the gates differ by condition because the declared architectures differ. The common cross-condition text outcome is final-output exactness, which asks only whether the complete user-visible output equals the requested text. Both are reported so that a process failure is not mistaken for a textual mismatch, or vice versa. No weighted composite score was used. Terminal errors counted as failures in the intention-to-observe denominator.

We also report:

- exact and declared-normalized text match;
- answer presence and span coverage;
- strict and tolerant reference selection;
- tool invocation, requested lookup, and bypass;
- placeholder parsing and deterministic replacement; and
- final-output exactness and extraneous text.

Normalization covered limited surface variation and was never used to erase a translation or edition difference. After the declared handling of wrapper markup, the strict endpoint compared the final text byte for byte.

Estimands and uncertainty. Condition rates use all 2,160 scheduled observations per condition and are accompanied by raw counts. The prospectively specified primary contrast is the paired risk difference between deterministic rendering and native generation. We also report paired source-supplied and tool-mediated contrasts. Paired intervals come from a 10,000-resample target-cluster bootstrap with seed 5601. Clustering by target acknowledges repeated observations of the same passage while preserving the crossed cells inside each resample.

Model, prompt, edition, and passage-stratum analyses are descriptive interactions. The study was not designed to rank vendors, estimate population prevalence across Scripture, or infer training-set membership. Model routes and editions are a fixed evaluation panel; the bootstrap does not support generalization to untested routes or editions.

We additionally report target-level rate distributions and a post-hoc linear probability sensitivity model with condition, route, edition, prompt-family, passage-stratum, and epoch fixed effects and

target-cluster-robust standard errors using a t_{19} reference distribution. This model was not part of the prospective lock and is used only to test whether the primary fixed-panel conclusions depend on the unadjusted presentation. It does not convert the purposive target panel into a probability sample.

5 Results

Three findings organize the results. First, every source-connected condition delivered exact text substantially more often than generation from model memory. Second, connecting a source did not remove failure; it moved failure to a different stage of the source-delivery chain. Third, passage length and the need to infer a reference exposed weaknesses that pooled condition rates alone would miss. We confirm the design completed, then give the primary comparison and trace those failures through the chain.

5.1 The design completed as scheduled

All 8,640 scheduled observations were obtained. Thirty-three ended in terminal provider errors, almost all in native generation, and no failed observation was re-prompted or replaced. These count as failures in the primary denominator rather than being dropped, which is the conservative choice: excluding them raises native exactness only from 25.00% to 25.38% and tool-mediated exactness from 80.09% to 80.13%. Provider errors therefore do not explain any result below. Full execution provenance, including per-route error counts and disclosed metadata gaps, is in the replication package.

5.2 Source-connected systems quoted far more exactly than generation from memory




Figure 2: Common final-text exactness and architecture-adherent exactness by condition. The measures differ only for tool-mediated retrieval, where some exact final answers did not satisfy the declared tool-use gates.

Condition	Final exact	E2E-EXACT exact	N	E2E-EXACT rate
Native parametric quotation	540	540	2,160	25.00%
Source-supplied quotation	2,022	2,022	2,160	93.61%
Tool-mediated retrieval	1,778	1,730	2,160	80.09%
Deterministic rendering	1,813	1,813	2,160	83.94%

Table 1: Common final-text exactness and architecture-adherent exactness. All scheduled observations, including terminal provider errors, remain in the denominator.

Native parametric quotation produced an exact final quotation in one quarter of observations (Figure 2; Table 1); in other words, three out of four requests failed the strict endpoint when the model had to supply the passage from memory. Every source-connected intervention substantially improved on that baseline. The oracle condition had the highest pooled rate, followed by deterministic rendering and tool-mediated retrieval.

That ordering should not be read as a ranking of deployable architectures. The oracle condition is handed the correct span and therefore never fails at selection or access, while the other three must obtain the text themselves. Its rate is an upper bound on what perfect retrieval could achieve, not a competitor’s score. The clearest evidence is in the prompt-family results below: when requests became harder to resolve, every condition that performs selection degraded, and the oracle condition did not move (Table 7). Section 5 returns to this with a comparison that holds the source constant.

Treatment vs. native	Paired risk difference	Cluster bootstrap 95% interval
Source supplied	+68.61 percentage points	[+61.71, +75.60]
Tool mediated	+55.09 percentage points	[+46.48, +63.47]
Deterministic rendering	+58.94 percentage points	[+51.16, +66.67]

Table 2: Paired effects on architecture-adherent exactness. Complete pairs: 2,160 for each contrast.

The prospectively specified deterministic-versus-native risk difference was +58.94 percentage points (Table 2). The source-supplied contrast was larger at +68.61 percentage points, and the tool-mediated contrast was +55.09 percentage points. The intervals describe variation across the twenty target clusters, not a probability sample of all possible Bible passages or user requests.

Condition	Target median	IQR	Minimum	Maximum
Native parametric	18.98%	[10.65, 39.58]	0.00%	56.48%
Source supplied	98.61%	[92.82, 100.00]	55.56%	100.00%
Tool mediated	88.43%	[72.45, 94.44]	34.26%	96.30%
Deterministic rendering	89.35%	[83.33, 92.59]	46.30%	96.30%

Table 3: Distribution of architecture-adherent exactness across the twenty targets. Each target rate pools the balanced route, edition, prompt, and epoch panel ($N = 108$ per condition and target).

The target distributions show that the pooled effects were not driven by one or two passages (Table 3), but they also expose substantial target heterogeneity. In the post-hoc cluster-robust sensitivity model, adjusted risk differences were +68.61 percentage points for source supplied (t_{19} interval [+60.83, +76.39]), +55.09 percentage points for tool mediated ([+45.68, +64.50]), and +58.94 percentage points for deterministic rendering ([+50.36, +67.51]). These estimates are nearly identical to the paired fixed-panel contrasts and are interpreted as sensitivity results, not new

confirmatory claims.

5.3 Each architecture left a different model decision exposed

Condition	Answered	Exact text	Method adherence	Final exact	E2E-EXACT
Native	98.29%	25.14%	–	25.00%	25.00%
Source supplied	100.00%	93.66%	100.00%	93.61%	93.61%
Tool mediated	99.81%	82.31%	84.17%	82.31%	80.09%
Deterministic rendering	100.00%	83.98%	92.55%	83.94%	83.94%

Table 4: Selected component rates. “Method adherence” is condition-specific and should not be compared as a general model capability.

The aggregate rates become more informative when read as a sequence of gates (Table 4). Source-supplied quotation removed the first two gates: the model did not have to identify or retrieve the passage because the authoritative span was already in context. Even so, 138 final outputs were inexact. Placing authoritative text in context brought the system close to deterministic copying without achieving it.

The tool condition reintroduced source access as a model decision. A tool invocation was observed in 2,052 of 2,160 observations (95.00%), but the tool was called for the requested span in only 1,818 (84.17%). The remaining trace-level failures included 107 bypasses and 234 calls for a span other than the requested span, plus one terminal error. Among observations satisfying tool method adherence, 1,730 of 1,818 (95.16%) achieved architecture-adherent exactness. The tool itself was not the guarantee; correct use of the tool was.

The deterministic condition moved responsibility again. Strict reference selection was correct in 1,815 of 2,160 observations (84.03%). Among those correct selections, 1,813 produced an exact final quotation (99.89%). Once a correct reference reached the renderer, its output was almost always exact. The dominant uncertainty was upstream: whether the model selected the right passage.

Formatting and selection were related but not identical. The outer placeholder was recognized in 2,159 observations, while 161 observations carried a syntax-related failure tag, usually because the model added an edition label to the reference payload. The analysis-only tolerant parser identified the expected reference in 1,954 observations (90.46%), including 158 of those 161 syntax-tagged cases. This does not change the strict 83.94% primary result.

Separating the two by prompt family shows that they fail in opposite circumstances (Table 5). When the caller named the reference, the model conformed to the strict grammar in only 85.56% of observations, yet the tolerant parser recovered the intended reference in 99.91%: the model almost always knew which passage was wanted and simply wrote the placeholder wrongly. When the caller described the passage instead, the pattern reversed— grammar conformance rose to 99.54% while recovered references fell to 81.02%, because the difficulty had moved from formatting to identification.

Interface conformance and semantic reference selection are therefore different questions, and a single accuracy figure conflates them. The strict endpoint charges both to the same account, which is the conservative choice for a locked comparison but understates how often the model identified the right passage.

Deterministic condition	Explicit reference	Contextual description
Strict grammar satisfied	85.56%	99.54%
Intended reference recovered	99.91%	81.02%

Table 5: Interface conformance against reference identification, by caller prompt family ($N = 1,080$ per cell). Recovery uses the prospectively specified analysis-only tolerant parser and never altered the primary score.

5.4 Holding the source constant isolates the rendering stage

The pooled condition rates in Table 1 compare architectures that carry different burdens: the oracle condition is handed the correct span, while the others must obtain it. A comparison that holds the source constant asks a narrower and more actionable question—once a system has the right text or the right reference, how faithfully does it deliver it?

Stage reached	Exact	N	Text wrong
Exact span supplied in context	2,022	2,160	6.39%
Tool retrieved the requested span	1,730	1,818	4.84%
Correct reference reached the renderer	1,813	1,815	0.11%

Table 6: Delivery fidelity conditional on having obtained the correct source. Each row is restricted to observations that reached the corresponding stage, so selection and access failures are excluded from all three.

Table 6 isolates the stage that follows source acquisition. A generative layer handed the exact text still altered it in 6.39% of observations, and a generative layer that had just retrieved the text through an authorized tool altered it in 4.84%. A non-generative renderer given a correct reference altered it in 0.11%—roughly a fiftyfold reduction relative to supplying the text in context.

This is the strongest practical result in the study, and it is narrower than the pooled comparison. It says nothing about which architecture delivers correct text most often end to end, because it conditions away the stage where deterministic rendering actually fails. It says that if the correct source has been obtained, generation is the wrong mechanism for reproducing it.

Both readings are reported because they answer different questions. The locked primary endpoint asks whether the user received the requested text at all. Table 6 asks which component should own the words once the source is settled. Neither replaces the other, and the components combined here were among the outcomes specified before execution (Section 4.6).

5.5 Contextual requests mainly challenged reference identification

Condition	Explicit reference	Contextual description
Native	27.41%	22.59%
Source supplied	93.52%	93.70%
Tool mediated	90.19%	70.00%
Deterministic rendering	85.56%	82.31%

Table 7: Architecture-adherent exactness by caller prompt family ($N = 1,080$ per cell).

The explicit-reference column is the case where a user already knows which passage they want

and says so. It is the narrower question of delivery, and it removes identification as a variable for every architecture. Even there, generation from memory succeeded in only 27.41% of observations: naming the passage does not make a model reproduce it.

Contextual descriptions had little pooled effect on the oracle condition, where the authoritative span was already present—further evidence that this condition does not participate in selection. Their largest effect appeared in tool-mediated retrieval, where exactness fell by 20.19 percentage points relative to explicit references. Deterministic rendering fell by 3.25 percentage points, consistent with the model having to identify a span before the renderer could act.

5.6 The requested edition mattered most without a source

Condition	BSB	WEBU	LSV
Native	41.53%	28.47%	5.00%
Source supplied	94.44%	93.75%	92.64%
Tool mediated	80.97%	79.72%	79.58%
Deterministic rendering	87.92%	83.61%	80.28%

Table 8: Architecture-adherent exactness by requested edition ($N = 720$ per cell).

Native exactness varied sharply by edition (Table 8), reaching 41.53% for BSB and 5.00% for LSV. Source-connected conditions were less edition-sensitive. These results establish an interaction with requested edition in this panel. They do not identify its cause and cannot be used to infer undisclosed training-data frequency.

5.7 Long passages remained difficult even with source access




Figure 3: Architecture-adherent exactness by passage stratum ($N = 540$ per cell). “Indirect event” describes target construction; both prompt families were evaluated for every target.

Long passages were the most difficult stratum under every architecture (Figure 3). Native exactness fell to 5.37%. Even when the source was supplied, long-passage exactness was 85.93%, indicating that rendering failures were not limited to source selection. Tool and deterministic conditions also lost accuracy through span selection before exact text could be delivered.

5.8 Effects varied by route but persisted across repetitions

Every route improved substantially over its own native condition, but the relative ordering of interventions differed (Appendix Table 9). This heterogeneity cautions against treating “RAG,” “tool use,” or “deterministic rendering” as model-independent properties.

Aggregate epoch rates were stable: 24.17–25.56% for native, 93.47–93.75% for source supplied, 79.72–80.69% for tool mediated, and 83.47–84.17% for deterministic rendering. Cell-level behavior was not perfectly stable. Of 720 matched cells per condition, all three epochs succeeded in 122 native, 650 source-supplied, 504 tool-mediated, and 537 deterministic cells. Repetition therefore added information beyond a single run.

5.9 Failure review located errors at different stages

A route-blinded review of sampled failures, conducted after primary scores were frozen, shows what each architecture’s failures actually looked like.

Native failures included near-exact surface differences, lexical drift toward another familiar rendering, incomplete spans, blended wording, and refusals. Source-supplied failures included both typography-only differences and substantive alteration or omission. Tool failures separated into bypass, tool-like text without an executable call, and genuine calls that retrieved the wrong boundaries. Deterministic selection failures frequently expanded an event to include narrative setup, narrowed a requested range to its most salient verse, or appended an edition label that violated the strict interface.

The pattern is consistent with the staged reading above: each architecture failed at the decision it still owned. These are mechanisms, not prevalence estimates; the sampling procedure is documented in the replication package, and no raw response is reproduced here.

6 Discussion

6.1 Quotation reliability moved through the source-delivery chain

The simplest reading of the primary table is that source-connected architectures outperformed generation from memory. That is true, but incomplete. The more useful finding is that each architecture made a different part of the journey reliable.

The oracle condition removed identification and retrieval by construction, leaving preservation as the only risk—and preservation still failed 6.39% of the time. Tool mediation gave the model access to a source of record, but also gave it decisions about whether to call the tool and which span to request. Deterministic rendering made the words nearly certain after a successful handoff, while leaving reference selection exposed. The failures did not disappear as architecture changed; they migrated toward the remaining model-owned decisions.

Two request properties made that migration visible. Contextual descriptions barely mattered when the exact source was already supplied, but exposed boundary and selection errors whenever the system had to identify a passage first: this is not noise around the same task, it is an additional reasoning stage. And the requested edition mattered most when the model answered from memory, which is where edition identity is least constrained.

The study therefore does not show that deterministic rendering is the best architecture overall; its own selection stage succeeded only 84.03% of the time. What it shows is narrower and more useful: once the correct source has been obtained, a non-generative component reproduces it far more reliably than a generative one, by 0.11% against 6.39%. Guarantees must be assigned to stages. A renderer can guarantee the bytes it emits without guaranteeing that the model asked

for the right source. A tool can make authoritative text available without guaranteeing that the model will use it. And a source placed in context can still be changed by the generative layer that presents it.

6.2 Implications for frontier AI systems

For frontier labs and application developers, the results suggest five design principles.

1. **Treat exact quotation as an intent class.** A system should distinguish requests for meaning from requests for exact source text.
2. **Make the source of record explicit.** Work, edition, span, authorization, and retrieval provenance should be represented in the system state rather than inferred only from generated prose.
3. **Separate selection from rendering.** Structured selection makes upstream uncertainty observable and permits a non-generative component to own final text. Measure the interface separately from the choice it carries: when the caller named the reference, our strict grammar was satisfied 85.56% of the time while the intended reference was recoverable 99.91% of the time, so a single score would have charged a formatting defect to comprehension.
4. **Verify tool behavior, not tool availability.** A tool definition does not establish invocation, correct lookup, or faithful rendering.
5. **Test the final user-visible buffer.** Exact source bytes can still be lost through truncation, wrappers, streaming errors, or later generative transformation.

These principles follow from the chain, not from Scripture, and apply wherever a source of record exists. The measured effect sizes do not: those require domain-specific replication.

6.3 Implications for Christian institutions

The results give Christian organizations a reason to distinguish “the model knows this passage” from “the system delivered this edition faithfully.” Native generation was especially unreliable for long passages and for the LSV edition. A fluent response should not be assumed to preserve the requested translation.

Three checks follow directly from the results and can be applied to a system an institution is evaluating. Ask whether the requested edition is named and carried through to the delivered text rather than inferred. Ask for evidence that an authorized source was actually consulted for the requested span, since the availability of a lookup tool did not establish its correct use here. And test long passages and event-style requests specifically, because both degraded under every architecture we evaluated.

These results also put a measured figure beside a widely circulated industry estimate that models misquote Scripture between 15% and 60% of the time ([Christian Daily International, 2026](#)). Native generation failed our strict endpoint in 75.00% of observations, which falls inside that range without replicating it: our endpoint demands an exact match to a specified edition on a purposive twenty-passages panel, and the circulated estimate states no protocol. The comparison is offered as calibration, not confirmation.

At the same time, textual fidelity should not be confused with theological fidelity. An exact verse can be detached from its literary context, used to answer the wrong question, or interpreted contrary to a community’s confessional commitments. The present evaluation can help institutions

remove misquotation as one failure mode. It cannot decide which passage should be used, how it should be interpreted, or what pastoral counsel should follow. Those tasks require human judgment and, in many settings, accountable ecclesial authority.

7 Open Evaluation and Reproducibility

The public repository, <https://github.com/FideAI/scripture-quotation-fidelity>, releases the deidentified 8,640-row derived-score dataset, the scoring specification and schemas, the executed prompts and tool schema, the source-edition record with per-target text digests, the prospective-lock receipt, and standalone analysis code. Together these reproduce every quantitative claim in this paper without network access or model spend.

Three things are withheld: authoritative passage text, which carries rights obligations; raw model outputs and provider response identifiers, which carry custody obligations; and private execution code and credentials. Withholding raw outputs is a real limit—a reader cannot inspect what a model actually said—and it is the reason the failure review reports mechanisms rather than transcripts. An exact rerun against live endpoints is in any case not guaranteed, since routing, model versions, and provider defaults change.

The shared condition implementation was executed from a pinned revision of The Apologist Project’s `llm-scripture-fidelity` repository (The Apologist Project, 2026). Fide AI controlled the crossed design, credentials, execution, scoring review, claims boundary, and result custody. The collaboration does not constitute product certification or commercial endorsement. The partner may publish a separate companion explanation of its implementation.

8 Responsible Use and Limitations

This work is a measurement study. It is not legal advice, copyright clearance, theological certification, pastoral-safety approval, product endorsement, or deployment authorization.

Rights status constrains which architectures are available at all: an edition that cannot be sent to a model cannot be supplied in context or returned by a tool. BSB and WEBU were evaluated as public-domain sources; LSV was used under its stated open license, with no LSV passage text included in the public result package. Restricted translations require separate source and source-to-model authorization. Technical ability to reproduce text does not itself establish a right to distribute it.

Throughout this paper, “fidelity” has a deliberately narrow meaning: fidelity to the requested English source text and span. Christian traditions also use the word for faithfulness to Scripture’s meaning, canonical context, doctrine, and lived obedience. Nothing in the experiment reduces those questions to string matching.

8.1 Limitations of the current study

Purposive target panel. The twenty passages were selected to expose distinct technical behaviors. They do not estimate the prevalence of quotation failure across the Bible, all Christian use cases, or ordinary user traffic.

English and Scripture only. The experiment used three English editions and one sacred-text domain. Reference conventions, translation ecosystems, source authority, and language technology differ elsewhere. Cross-lingual and cross-domain claims require new studies.

The oracle condition is not a deployable comparison. Source-supplied quotation is handed the correct span and therefore bypasses selection and access entirely. Its rate bounds what perfect retrieval could deliver; it does not estimate what any retrieval-augmented system would achieve, since a real system must decide what to fetch. Comparisons that place it alongside the three deployable architectures overstate it by exactly the cost of source selection, which for the deterministic condition was the difference between 84.03% and 99.89%. A study that gave every condition the same selection burden—for instance by having the model choose a reference, then retrieving it, then supplying it—would answer a different and also useful question, and would require its own prospective specification.

Edition versioning. Ground truth was a fixed stored copy of each edition, verified stable by repeat-fetch digests across the run and disclosed in the released source-edition record. Open English editions are nonetheless revised over time, so an exact replication requires the same edition version, not merely the same edition name; the released per-target digests let a replicator confirm that condition before running anything. The digests establish that our ground truth was internally stable and externally checkable, not that the upstream provider’s text is itself authoritative for a given publisher.

Open editions only. The primary experiment does not measure restricted translations. It therefore cannot establish how licensing, provider quotation policies, or source-transfer permissions affect the four architectures.

Model and route drift. Results apply to six named routes on July 26–27, 2026. Hosted models, routing, guardrails, and provider defaults can change. Three epochs measure short-run variability, not longitudinal stability.

Provider-default sampling. Temperature and other generation controls were omitted and left at provider defaults. This keeps the comparison close to how these endpoints behave in ordinary use, but it prevents any causal claim about sampling parameters.

Strict interface contract. The deterministic primary result includes failures caused by a deliberately strict reference grammar, and that cost is concentrated where the caller named the reference: grammar conformance there was 85.56% against 99.91% reference recovery. A production system would likely adopt a more permissive parser and would recover much of that gap. We do not report what such a system would have scored, because selecting a parser after observing results would invalidate the locked comparison; the two components are reported separately so a reader can see the size of the effect without our having chosen it.

Automated exactness. Exact matching is appropriate for the declared endpoint, but failure tags and normalized views depend on implementation choices. A post-run route-blinded review examined mechanisms in a small diagnostic sample; it did not re-annotate the full set of observations.

Dependence and fixed-panel inference. Target-cluster bootstrap intervals describe variation within this target panel, while model routes and editions remain fixed. Target-level distributions and a post-hoc target-cluster-robust fixed-panel sensitivity analysis make this dependence visible, but neither supports population inference to all Scripture requests, models, or editions. With only

twenty target clusters, uncertainty estimates remain coarse; population claims would require a larger probability sample of passages.

Contextual-description review. The six revised contextual descriptions received a structurally blinded technical reapproval, but the reviewer was an AI system and the same reviewing identity that authored the initial critique. This supports an internal construct check, not independent credentialed scholarly validation. The contextual-prompt results should be read with that limitation attached: a low reference-identification rate could reflect system behavior or description quality, and the present evidence does not fully separate them. Materials for an independent scholarly review of all twenty descriptions—protocol, blinded packet, and pre-committed scoring rules—are released with this package so that check can be carried out and reported by us or by others. Primary scores are frozen, so such a review can qualify how the contextual-prompt findings are interpreted but cannot alter any reported quantity.

No contextual or theological endpoint. Contextual descriptions test technical reference identification. The study does not determine whether the expected passage is the best theological, exegetical, or pastoral response to a real person’s situation.

9 Conclusion

AI systems should not always generate. In this study, native parametric quotation delivered the requested edition exactly in 25.00% of observations. Supplying authoritative text, connecting a source tool, or delegating final rendering to a deterministic layer all produced large improvements, but none eliminated failure. Each architecture moved failure to the decision it still left to the model.

The practical lesson is not that one method wins every comparison. It is that authoritative quotation should be designed and evaluated as a sequence of observable responsibilities. A model can identify what a user means. A source interface can retrieve what is authorized. A deterministic component can render what must not change. And a final-output check can establish what the user actually received.

Return to the opening request about Jesus and the serpent in the wilderness. The system should be able to show whether the model identified the intended passage, whether an authorized source supplied the requested edition, and whether the final words survived intact. If the answer is wrong, the same chain should reveal whether the error was selection, access, or rendering.

Scripture makes these distinctions visible, but the underlying problem belongs to any AI system that handles authoritative text. When wording is material, the system should be able to say when the model is speaking and when the source of record is speaking.

Data and Code Availability

The public repository, <https://github.com/FideAI/scripture-quotation-fidelity>, releases the full deidentified derived-score dataset, target registry, executed prompt templates, functional tool schema, source-edition record with per-target text digests, release-safe prospective-lock and deviation records, and standalone analysis code. It excludes generated prose, authoritative passage text, provider response identifiers, raw tool traces, credentials, restricted-source payloads, and private execution-engine code. The released package reproduces the quantitative tables and contrasts without model endpoint access.

Acknowledgments

We thank The Apologist Project, which developed and maintains the shared condition implementation used to execute the four architectures ([The Apologist Project, 2026](#)), and which contributed to how those conditions were operationalized: what each architecture had to do in practice for a quotation to count as delivered through it. That work shaped the instrument this study depends on, and the study would have been materially harder to run without it.

The Apologist Project is not an author of this paper. Fide AI set the research question and crossed design, held credentials, ran confirmatory execution, performed the analysis, held evidence custody, and made the release decision. Responsibility for the claims made here, and for any error in them, rests with the author alone.

Competing Interests

The Apologist Project develops commercial technology related to Scripture delivery and therefore holds an organizational interest relevant to this study. Because it also contributed to how the conditions were operationalized, that interest touches the instrument and not only the topic, and readers should weigh the results accordingly.

Three features of the study bear on that concern. Condition definitions, scoring rules, and the primary endpoint were fixed in an internal hash-lock before confirmatory execution. Fide AI, not the partner, held the credentials, executed the study, performed the analysis, and controlled the claims boundary. And the architecture closest to a deterministic Scripture-delivery product did not finish first: source-supplied quotation reached 93.61% against 83.94% for deterministic rendering, and the Discussion states that the study does not show deterministic rendering to be universally preferable.

The study did not evaluate or certify a deployed commercial product.

Use of AI Systems

AI systems assisted with code generation, manuscript drafting, citation cross-checking, and internal contextual-description review. Quantitative outcomes were produced by the declared automated scoring pipeline. The AI review is disclosed as an internal construct check and is not represented as independent credentialed scholarly validation.

A Target References

The twenty target references were: John 3:16; Genesis 1:1; Romans 8:28; Hebrews 11:1; John 14:6; Psalm 23:1–3; Proverbs 3:5–6; Ephesians 2:8–9; Philippians 4:6–7; Deuteronomy 6:4–5; Matthew 5:3–10; 1 Corinthians 13:4–7; Isaiah 53:4–6; Revelation 21:1–4; Matthew 22:37–40; Numbers 21:8–9; John 3:14–15; Luke 2:10–12; Micah 6:8; and James 1:5. The repository includes the exact contextual descriptions used for each.

B Execution Routes and Prospective Lock

The executed OpenRouter model identifiers and pinned upstreams were: `openai/gpt-5.6-sol` via OpenAI; `anthropic/claude-sonnet-5` via Anthropic; `google/gemini-3.5-flash` via Google

AI Studio; deepseek/deepseek-v4-pro via Together; moonshotai/kimi-k3 via Moonshot; and z-ai/glm-5.2 via Together. Fallbacks were disabled, required parameter support was enabled, and data-collecting routes were denied.

The internal prospective lock was recorded at 2026-07-26T22:28:40Z. Confirmatory execution began at 2026-07-26T23:11:30Z. The lock fixed the target pack, execution matrix, source registry, scoring rubric, and deviation log by SHA-256 digest.

C Route-by-Condition Results

Evaluated route	Native	Source supplied	Tool	Deterministic
Claude Sonnet 5	12.78%	82.22%	78.06%	75.56%
DeepSeek V4 Pro	32.22%	96.11%	83.61%	89.17%
Gemini 3.5 Flash	16.11%	98.06%	86.94%	91.94%
GLM-5.2	20.00%	93.61%	75.56%	85.83%
GPT-5.6 Sol	35.28%	97.22%	92.50%	71.39%
Kimi K3	33.61%	94.44%	63.89%	89.72%

Table 9: Descriptive route-by-condition architecture-adherent exactness ($N = 360$ per cell). Rows are ordered alphabetically, not by performance; this is not a model leaderboard.

D Failure Tags

Observed structured tags included 345 wrong or unparseable deterministic reference selections, 234 tool calls that did not retrieve the requested span, 161 invalid deterministic reference payloads or placeholder failures, 161 selected-reference lookup failures, and 107 tool-bypass observations. Tags can co-occur and must not be summed as unique failed observations. Native and source-supplied text mismatches are represented primarily through exact, normalized, coverage, and answer-presence metrics rather than these architecture-specific tags.

E Claims Boundary

Results apply only to the named routes, dates, twenty-target panel, two prompt families, three open English editions, four architecture conditions, and declared scoring protocol. The study does not establish:

- a general ranking of models or vendors;
- theological truth or contextual appropriateness;
- pastoral safety or institutional deployment readiness;
- training-data membership or memorization causation;
- universal legal or licensing compliance;
- performance on restricted translations; or
- superiority of a commercial product.

References

- Barna Group. 2026. New research: Pastors are using AI more than you think. <https://www.barna.com/research/pastors-using-ai-ministry/>. Published June 15, 2026. Online surveys of U.S. Protestant pastors ($n = 442$, December 2025) and U.S. adults ($n = 1,514$, November 2025); accessed 2026-07-29.
- Berean Bible. 2023. Free licensing. <https://berean.bible/licensing.htm>. Licensing statement; accessed 2026-07-27.
- Nicholas Carlini, Daphne Ippolito, Matthew Jagielski, Katherine Lee, Florian Tramèr, and Chiyuan Zhang. 2023. [Quantifying memorization across neural language models](#). In *The Eleventh International Conference on Learning Representations*. OpenReview.net.
- Nicholas Carlini, Florian Tramèr, Eric Wallace, Matthew Jagielski, Ariel Herbert-Voss, Katherine Lee, Adam Roberts, Tom Brown, Dawn Song, Úlfar Erlingsson, Alina Oprea, and Colin Raffel. 2021. [Extracting training data from large language models](#). In *30th USENIX Security Symposium (USENIX Security 21)*, pages 2633–2650. USENIX Association.
- Tong Chen, Akari Asai, Niloofar Mireshghallah, Sewon Min, James Grimmermann, Yejin Choi, Hannaneh Hajishirzi, Luke Zettlemoyer, and Pang Wei Koh. 2024. [CopyBench: Measuring literal and non-literal reproduction of copyright-protected text in language model generation](#). In *Proceedings of the 2024 Conference on Empirical Methods in Natural Language Processing*, pages 15134–15158, Miami, Florida, USA. Association for Computational Linguistics.
- Christian Daily International. 2026. AI’s bible misquotes range from 15% to 60%, YouVersion CEO says. <https://www.christianpost.com/news/ais-bible-misquotes-range-from-15-to-60-youversion-ceo.html>. Published March 17, 2026 in The Christian Post. Reported industry estimate; no measurement methodology given; accessed 2026-07-29.
- Covenant Press. 2020. Literal standard version. <https://www.bible.com/en-GB/versions/2660-lsv-literal-standard-version>. Edition and CC BY-SA licensing statement; accessed 2026-07-27.
- eBible.org. 2024. The world english bible (WEB) faq. <https://ebible.org/eng-web/webfaq.htm>. Updated February 5, 2024; accessed 2026-07-27.
- Shahul Es, Jithin James, Luis Espinosa Anke, and Steven Schockaert. 2024. [RAGAs: Automated evaluation of retrieval augmented generation](#). In *Proceedings of the 18th Conference of the European Chapter of the Association for Computational Linguistics: System Demonstrations*, pages 150–158, St. Julians, Malta. Association for Computational Linguistics.
- Exponential AI NEXT and ChurchTechToday. 2025. The 2025 state of AI in the church national survey report. <https://churchtechtoday.com/2025-state-of-ai-in-the-church/>. Survey of 594 church leaders; accessed 2026-07-29.
- Fide AI. 2026. FID-056: Formal verification for sacred text fidelity. <https://github.com/FideAI/research-ideas/blob/main/ideas/FID-056-formal-verification-sacred-text-fidelity.md>. Public research call; published at <https://fideai.org/research/calls>; accessed 2026-07-30.

- Tianyu Gao, Howard Yen, Jiatong Yu, and Danqi Chen. 2023. [Enabling large language models to generate text with citations](#). In *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing*, pages 6465–6488, Singapore. Association for Computational Linguistics. Introduces the ALCE benchmark.
- Saibo Geng, Martin Josifoski, Maxime Peyrard, and Robert West. 2023. [Grammar-constrained decoding for structured NLP tasks without finetuning](#). In *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing*, pages 10932–10952, Singapore. Association for Computational Linguistics.
- Philip Goff, Arthur E. Farnsley, II, and Peter J. Thuesen, editors. 2017. *The Bible in American Life*. Oxford University Press, New York. Includes General Social Survey and National Congregations Study modules on which translation respondents read.
- Lei Huang, Weijiang Yu, Weitao Ma, Weihong Zhong, Zhangyin Feng, Haotian Wang, Qianglong Chen, Weihua Peng, Xiaocheng Feng, Bing Qin, and Ting Liu. 2025. [A survey on hallucination in large language models: Principles, taxonomy, challenges, and open questions](#). *ACM Transactions on Information Systems*, 43(2):1–55.
- Ziwei Ji, Nayeon Lee, Rita Frieske, Tiezheng Yu, Dan Su, Yan Xu, Etsuko Ishii, Ye Jin Bang, Andrea Madotto, and Pascale Fung. 2023. [Survey of hallucination in natural language generation](#). *ACM Computing Surveys*, 55(12):1–38.
- Patrick Lewis, Ethan Perez, Aleksandra Piktus, Fabio Petroni, Vladimir Karpukhin, Naman Goyal, Heinrich Küttler, Mike Lewis, Wen-tau Yih, Tim Rocktäschel, Sebastian Riedel, and Douwe Kiela. 2020. [Retrieval-augmented generation for knowledge-intensive nlp tasks](#). *Advances in Neural Information Processing Systems*, 33:9459–9474.
- Percy Liang, Rishi Bommasani, Tony Lee, Dimitris Tsipras, Dilara Soylu, Michihiro Yasunaga, Yian Zhang, Deepak Narayanan, Yuhuai Wu, Ananya Kumar, Benjamin Newman, Binhang Yuan, Bobby Yan, Ce Zhang, Christopher Cosgrove, Christopher D. Manning, Christopher Ré, Diana Acosta-Navas, Drew A. Hudson, and 31 others. 2023. [Holistic evaluation of language models](#). *Transactions on Machine Learning Research*.
- Alex Mallen, Akari Asai, Victor Zhong, Rajarshi Das, Daniel Khashabi, and Hannaneh Hajishirzi. 2023. [When not to trust language models: Investigating effectiveness of parametric and non-parametric memories](#). In *Proceedings of the 61st Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 9802–9822, Toronto, Canada. Association for Computational Linguistics.
- Bruce M. Metzger and Bart D. Ehrman. 2005. *The Text of the New Testament: Its Transmission, Corruption, and Restoration*, 4 edition. Oxford University Press, New York.
- Sewon Min, Kalpesh Krishna, Xinxi Lyu, Mike Lewis, Wen-tau Yih, Pang Koh, Mohit Iyyer, Luke Zettlemoyer, and Hannaneh Hajishirzi. 2023. [FActScore: Fine-grained atomic evaluation of factual precision in long form text generation](#). In *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing*, pages 12076–12100, Singapore. Association for Computational Linguistics.
- Eugene A. Nida and Charles R. Taber. 1969. *The Theory and Practice of Translation*, volume 8 of *Helps for Translators*. E. J. Brill, Leiden.

- Marco Tulio Ribeiro, Tongshuang Wu, Carlos Guestrin, and Sameer Singh. 2020. [Beyond accuracy: Behavioral testing of NLP models with CheckList](#). In *Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics*, pages 4902–4912, Online. Association for Computational Linguistics.
- Hayley Ross, Ameya Sunil Mahabaleshwarkar, and Yoshi Suhara. 2025. [When2Call: When \(not\) to call tools](#). In *Proceedings of the 2025 Conference of the Nations of the Americas Chapter of the Association for Computational Linguistics: Human Language Technologies (Volume 1: Long Papers)*, pages 3391–3409, Albuquerque, New Mexico. Association for Computational Linguistics.
- Timo Schick, Jane Dwivedi-Yu, Roberto Dessì, Roberta Raileanu, Maria Lomeli, Eric Hambro, Luke Zettlemoyer, Nicola Cancedda, and Thomas Scialom. 2023. [Toolformer: Language models can teach themselves to use tools](#). In *Advances in Neural Information Processing Systems*, volume 36.
- The Apologist Project. 2026. [llm-scripture-fidelity: shared condition implementation for scripture quotation evaluation](#). Pinned revision executed for this study. Archived at Software Heritage: [swh:1:rev:b67849d55181fdffb3ee28d405733dd5eb35e4bc](#); accessed 2026-07-31.
- Emanuel Tov. 2012. *Textual Criticism of the Hebrew Bible*, 3, revised and expanded edition. Fortress Press, Minneapolis.
- United States Copyright Office. 2025. Copyright and artificial intelligence, part 3: Generative ai training (pre-publication version). <https://www.copyright.gov/ai/Copyright-and-Artificial-Intelligence-Part-3-Generative-AI-Training-Report-Pre-Publication-Version.pdf>. Report of the Register of Copyrights, released May 2025.
- Jingyu Zhang, Marc Marone, Tianjian Li, Benjamin Van Durme, and Daniel Khashabi. 2025. [Verifiable by design: Aligning language models to quote from pre-training data](#). In *Proceedings of the 2025 Conference of the Nations of the Americas Chapter of the Association for Computational Linguistics: Human Language Technologies (Volume 1: Long Papers)*, pages 3748–3768, Albuquerque, New Mexico. Association for Computational Linguistics.